

# Timeline

60min [09:00 - 10:00] **Intro & Fundamentals**

15min [10:00 - 10:15] Break

60min [10:15 - 11:15] **Supervised learning**

15min [11:15 - 11:30] Break

30min [11:30 - 12:00] **Unsupervised learning**

90min [12:00 - 13:30] Lunch

45min [13:30 - 14:15] **Pytorch & Keras overview**

15min [14:15 - 14:30] **Visual clutter classifier**

15min [14:30 - 14:45] **Fitts law regressor**

15min [14:45 - 15:00] Break

30min [15:00 - 15:30] **Clustering of UI wireframes**



SAARBRÜCKEN

# 6TH SUMMER SCHOOL ON COMPUTATIONAL INTERACTION

INFERENCE, OPTIMIZATION AND MODELING FOR THE  
ENGINEERING OF INTERACTIVE SYSTEMS | 13 - 18 JUNE 2022

## Deep Learning for Human–Computer Interaction **Session 1: Intro & Fundamentals**

**Luis Leiva & Bereket Yilma**  
University of Luxembourg



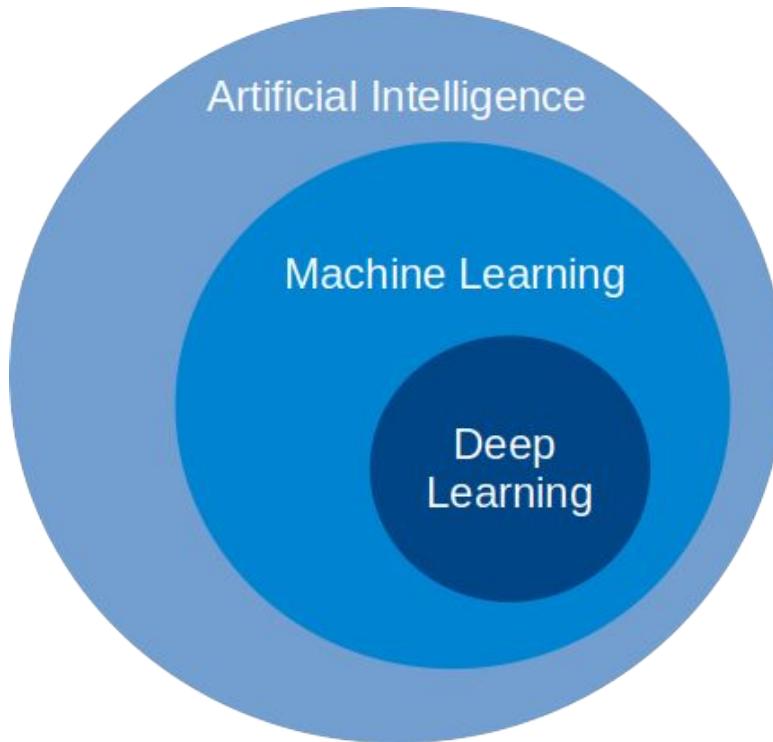
# Learning outcomes

After this lecture you will be able to:

- Understand ML and DL foundations
- Recognize key building blocks in DL models
- Formulate HCI problems from a ML/DL perspective

# Introduction

# Where we are



## Artificial Intelligence

Try to create machines that can simulate human behavior

## Machine Learning

Automatically learn from data without explicit programming

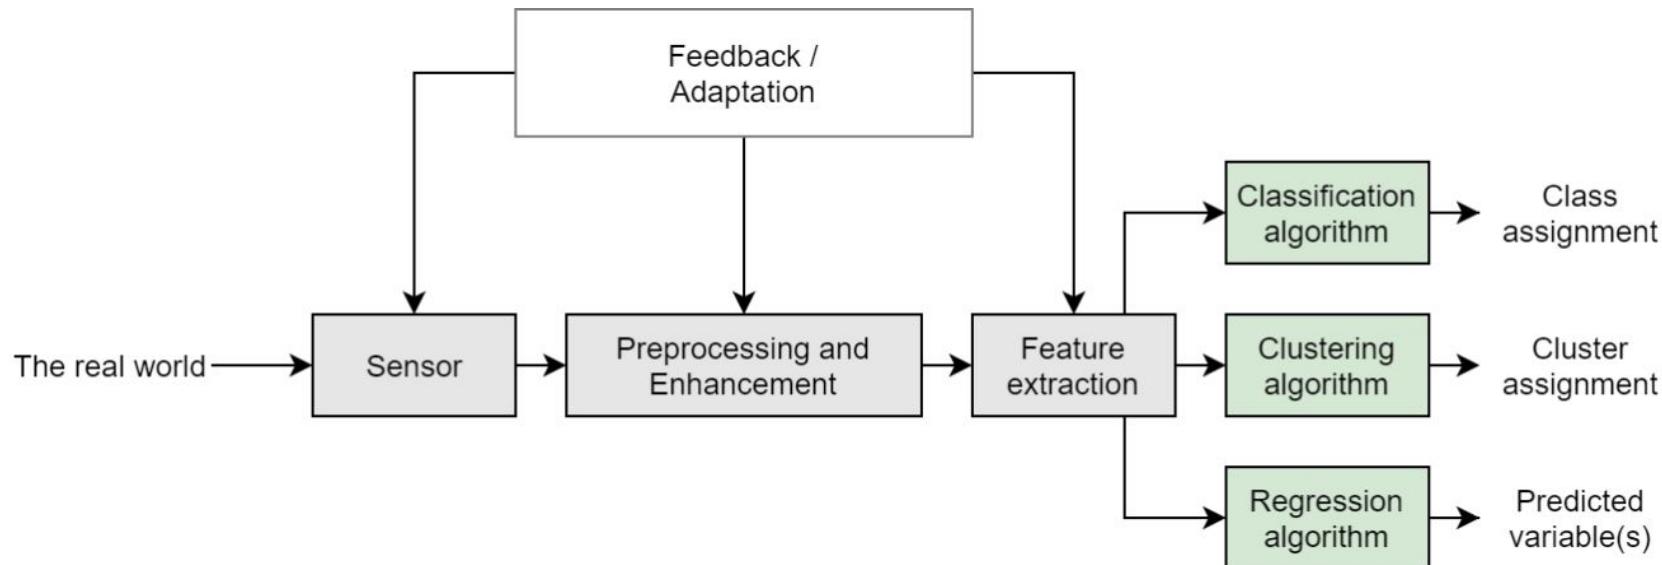
## Deep Learning

Machine learning with neural networks

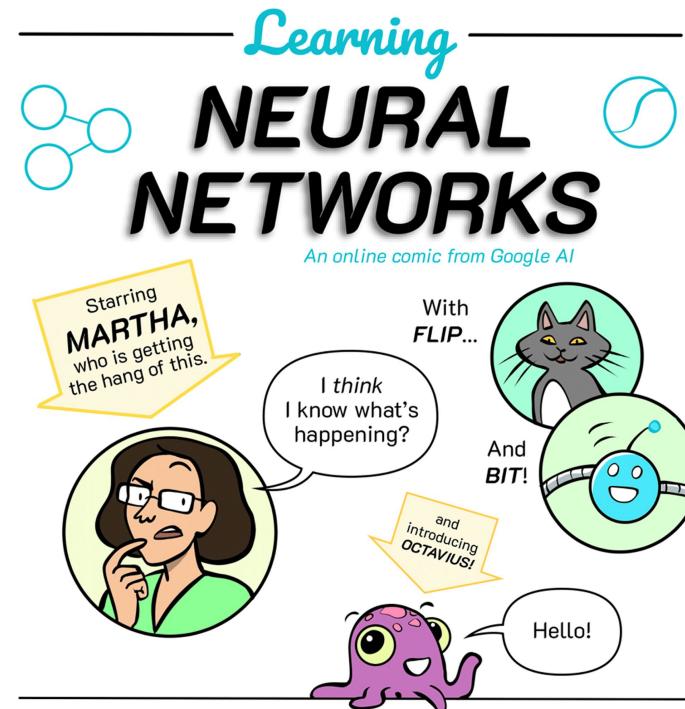
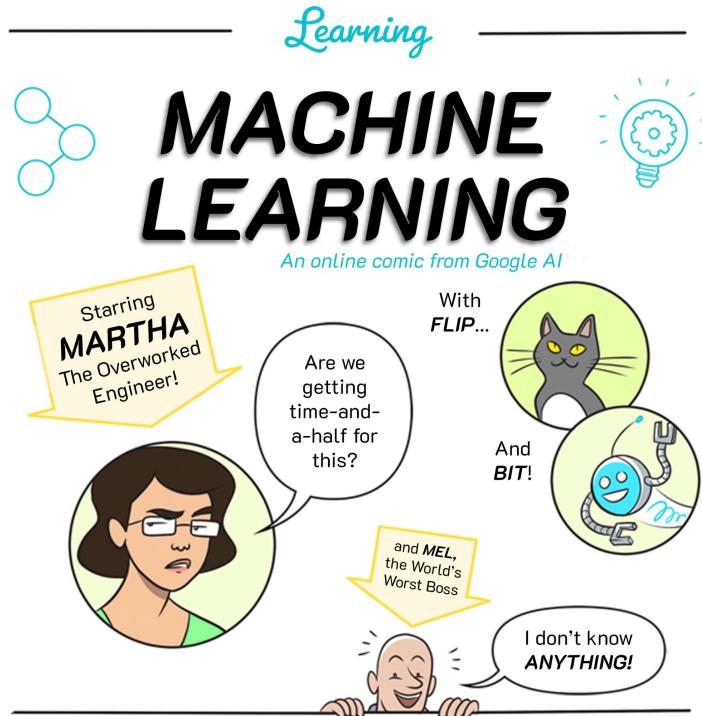
# Machine Learning

Pattern recognition and inference

Often guided by examples



# ML & DL preliminaries



# Why is it called DL?

Input layer

One or more hidden layers

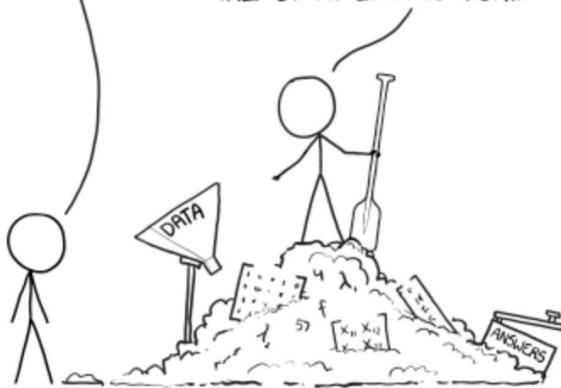
Output layer

THIS IS YOUR MACHINE LEARNING SYSTEM?

YUP! YOU POUR THE DATA INTO THIS BIG  
PILE OF LINEAR ALGEBRA, THEN COLLECT  
THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL  
THEY START LOOKING RIGHT.



<https://xkcd.com/1838/>

# Why DL now?



More **hardware**

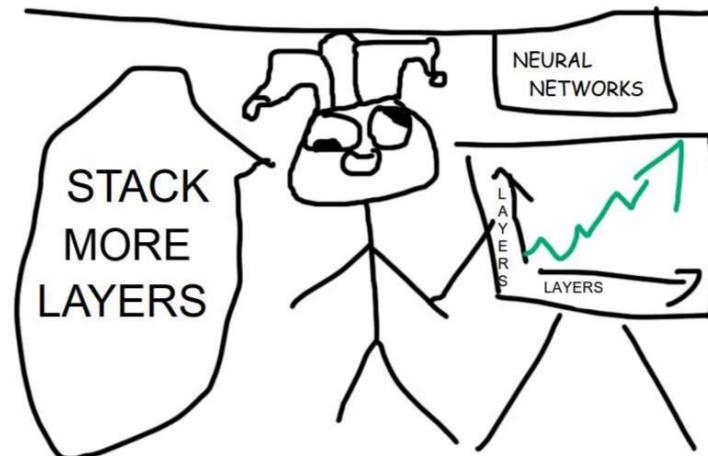
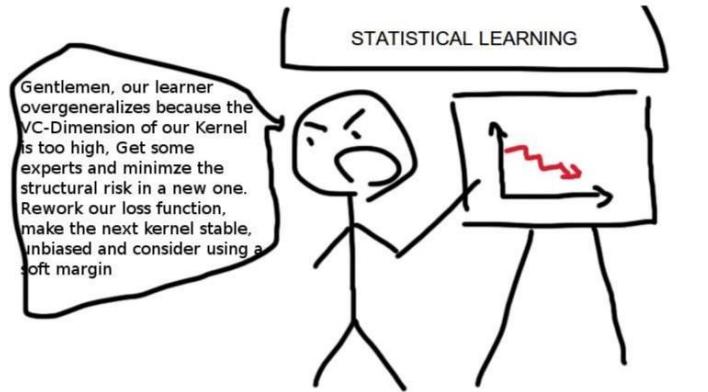


More **software**

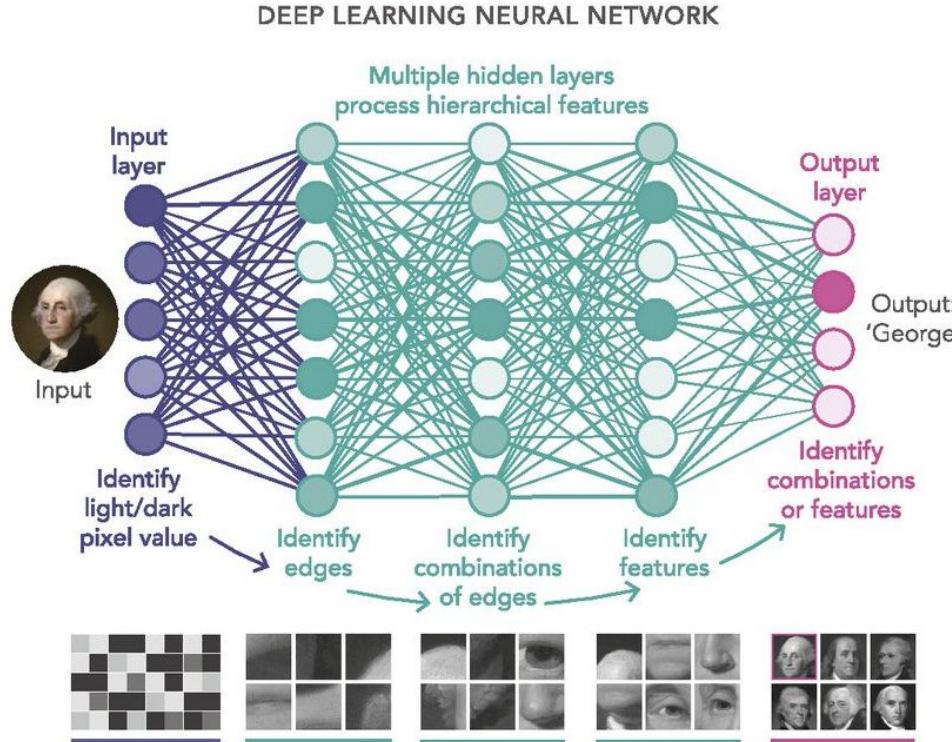


More **data**

# Why DL now?

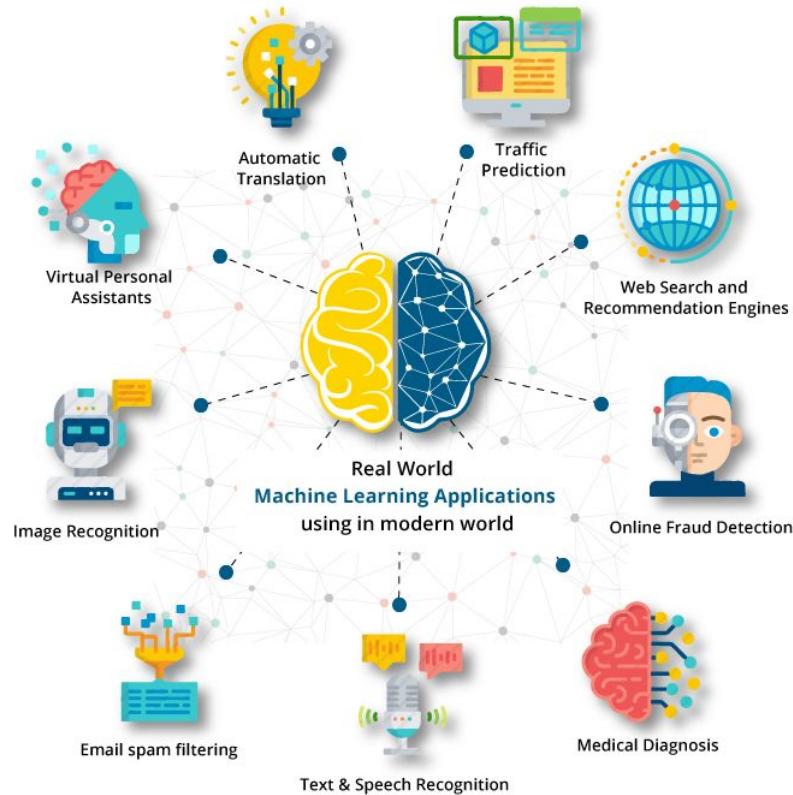
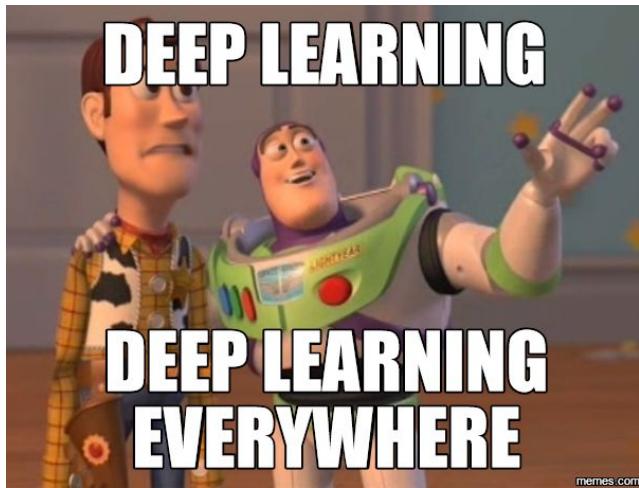


# Example



<https://www.pnas.org/content/116/4/1074>

# Applications



# Learning paradigms

## Supervised Learning

**Examples:**  
Classification  
Regression

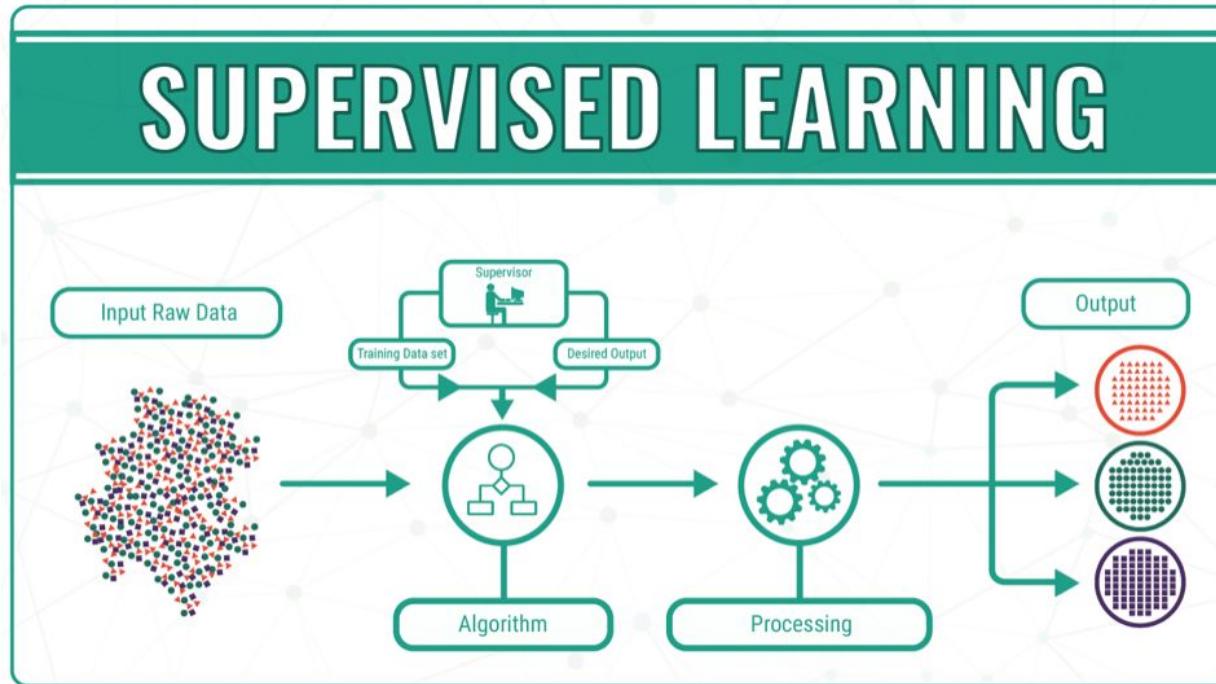
## Unsupervised Learning

**Examples:**  
Autoencoders  
Clustering

## Reinforcement Learning

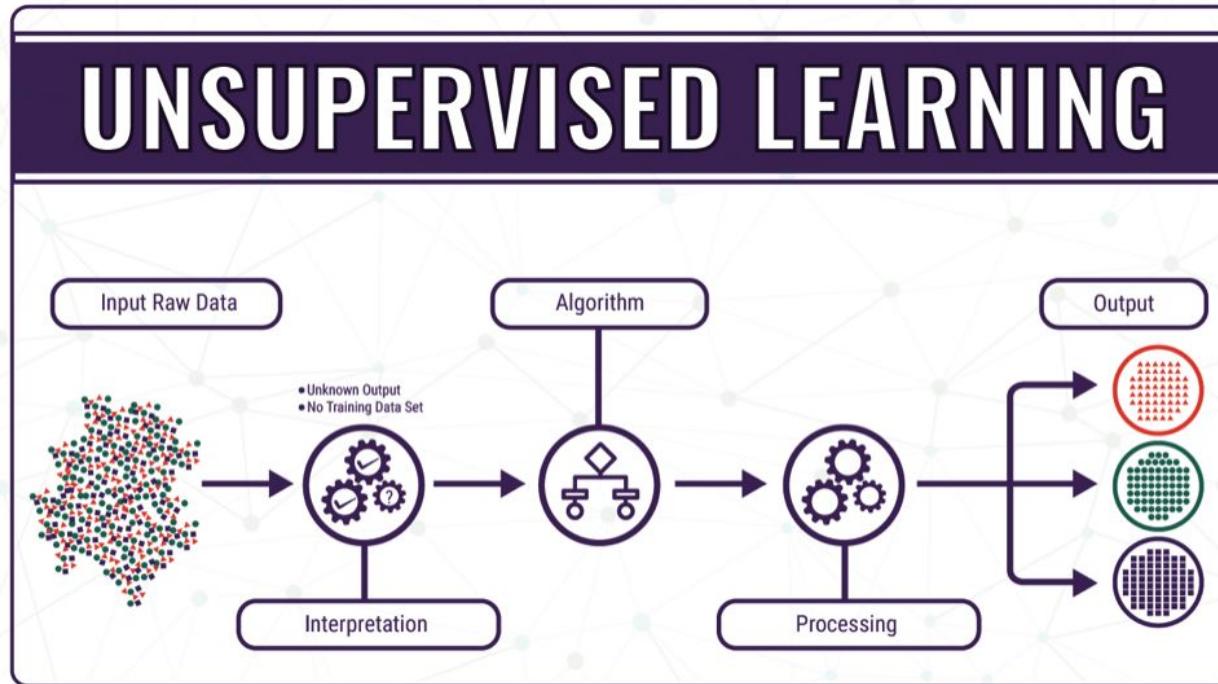
**Examples:**  
Navigation  
Planning

# Learning paradigms



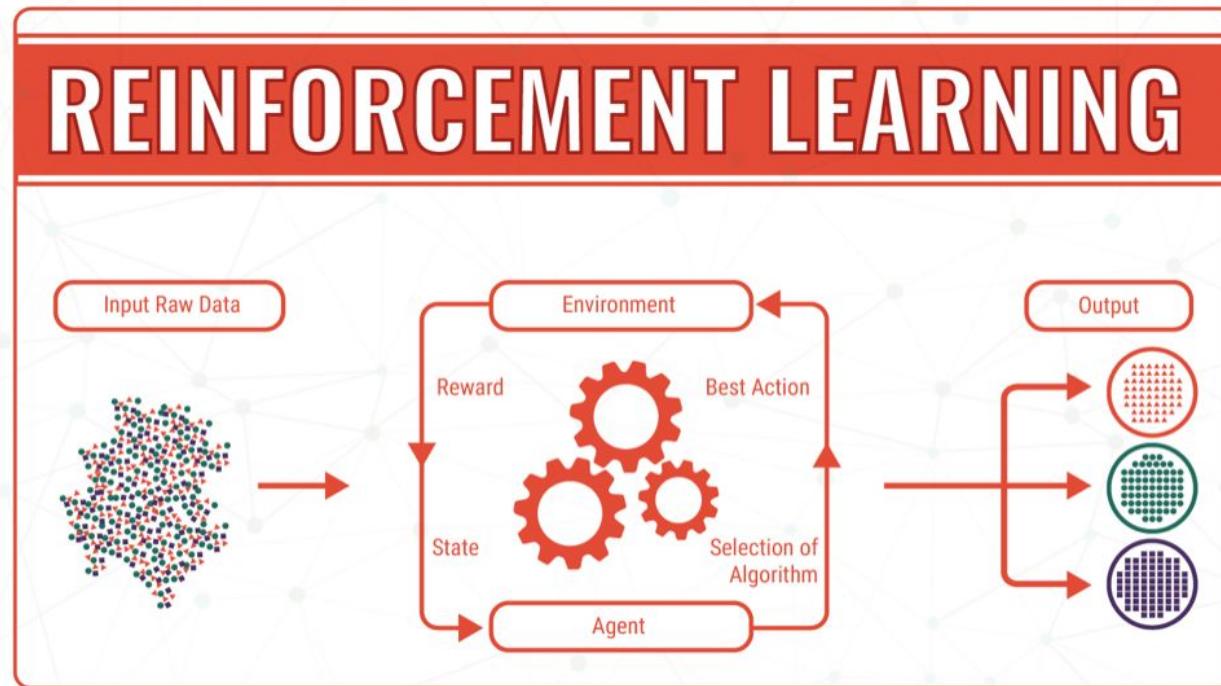
<https://www.linkedin.com/pulse/machine-learning-explained-understanding-supervised-ronald-van-loon>

# Learning paradigms



<https://www.linkedin.com/pulse/machine-learning-explained-understanding-supervised-ronald-van-loon>

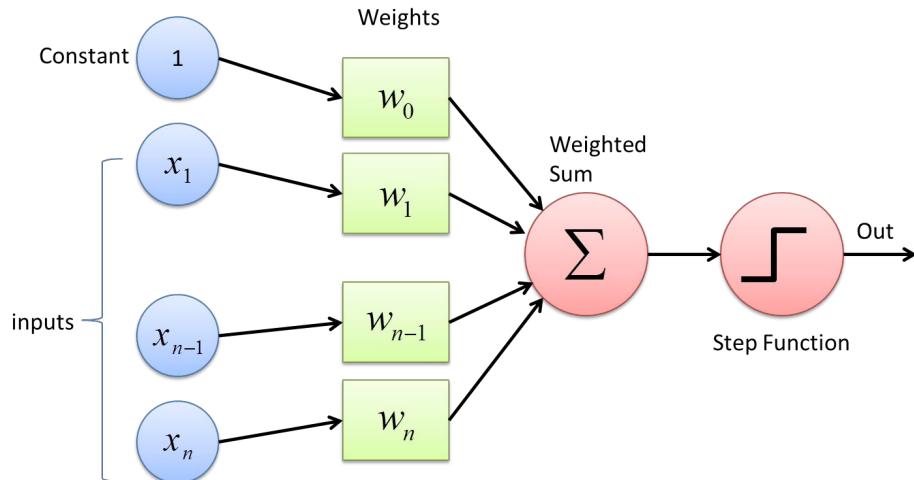
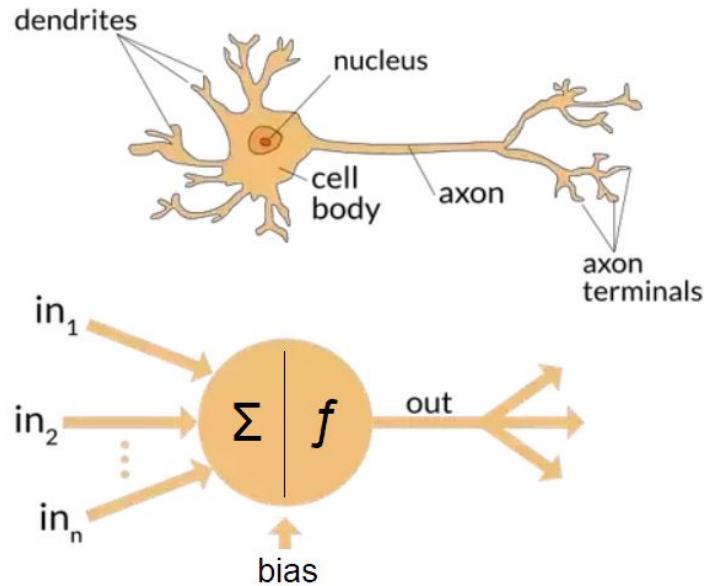
# Learning paradigms



<https://www.linkedin.com/pulse/machine-learning-explained-understanding-supervised-ronald-van-loon>

# Foundations

# Perceptron algorithm



<https://towardsdatascience.com/a8b46db828b7>

<https://towardsdatascience.com/626217814f53>

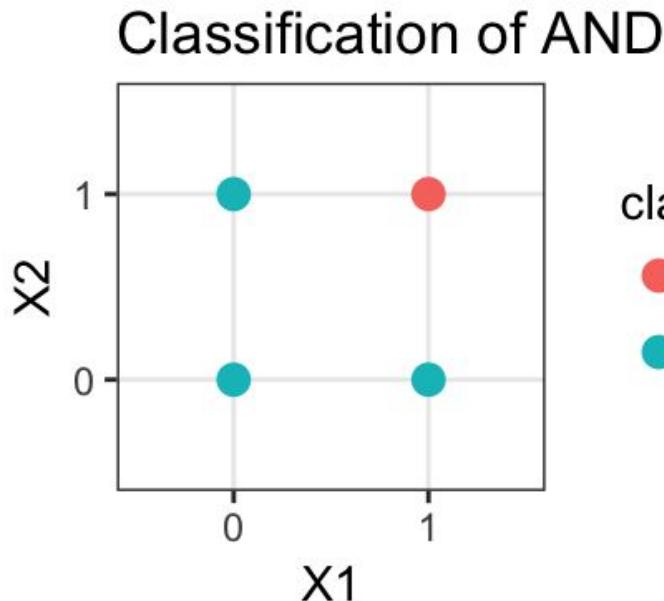
# Perceptron algorithm

1. Initialize weights  $\mathbf{w}$
2. For  $N$  iterations (or until convergence) do:
  - 2.1. Process each instance  $f(z) = f(\mathbf{w}^\top \mathbf{x})$
  - 2.2. Update weights:  $w_i += [y - f(z)] x_i$
3. Return weights

$$\hat{y} = f(z) = \begin{cases} 1 & w_0 + \sum_{i=1}^n x_i w_i > \theta \\ 0 & \text{otherwise} \end{cases}$$

$$\hat{y} = f(z) = \begin{cases} 1 & \sum_{i=0}^n x_i w_i > 0 \\ 0 & \text{otherwise} \end{cases}$$

# Perceptron algorithm



$x_1$	$x_2$	$y$
0	0	0
1	0	0
0	1	0
1	1	1

# Perceptron algorithm

Let  $y = f(b, x_1, x_2)$  and  $w = (0, 0, 0)$ .

Train  $x = (1, 1, 1)$ ;  $y = 1$

$$z = x \cdot w^T = 1 \cdot 0 + 1 \cdot 0 + 1 \cdot 0 = 0 \rightarrow f(z) = 0$$

$$w = (0, 0, 0) + (1 - 0) \cdot (1, 1, 1) = (1, 1, 1)$$

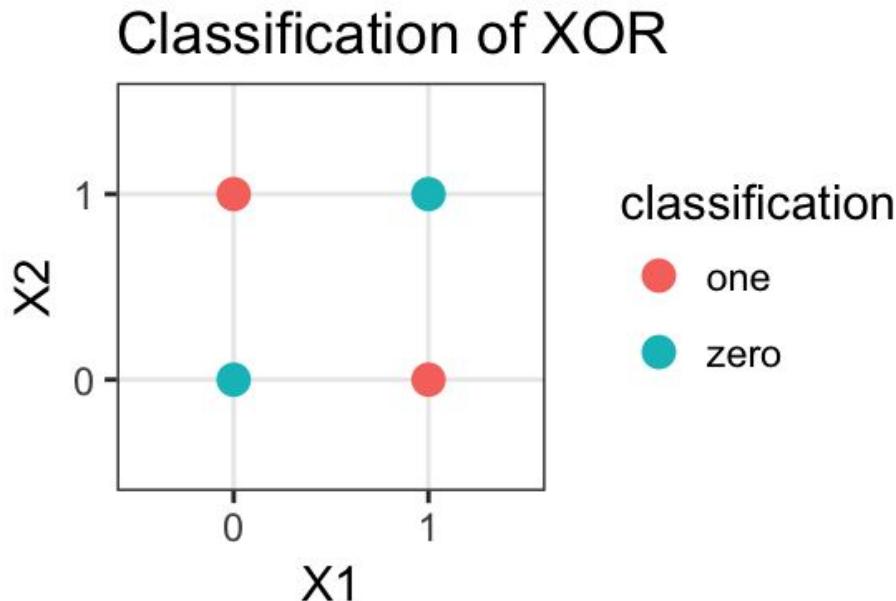
Train  $x = (1, 1, 0)$ ;  $y = 0$

$$z = x \cdot w^T = 1 \cdot 1 + 1 \cdot 1 + 0 \cdot 1 > 0 \rightarrow f(z) = 1$$

$$w = (1, 1, 1) + (0 - 1) \cdot (1, 1, 0) = (0, 0, 1)$$

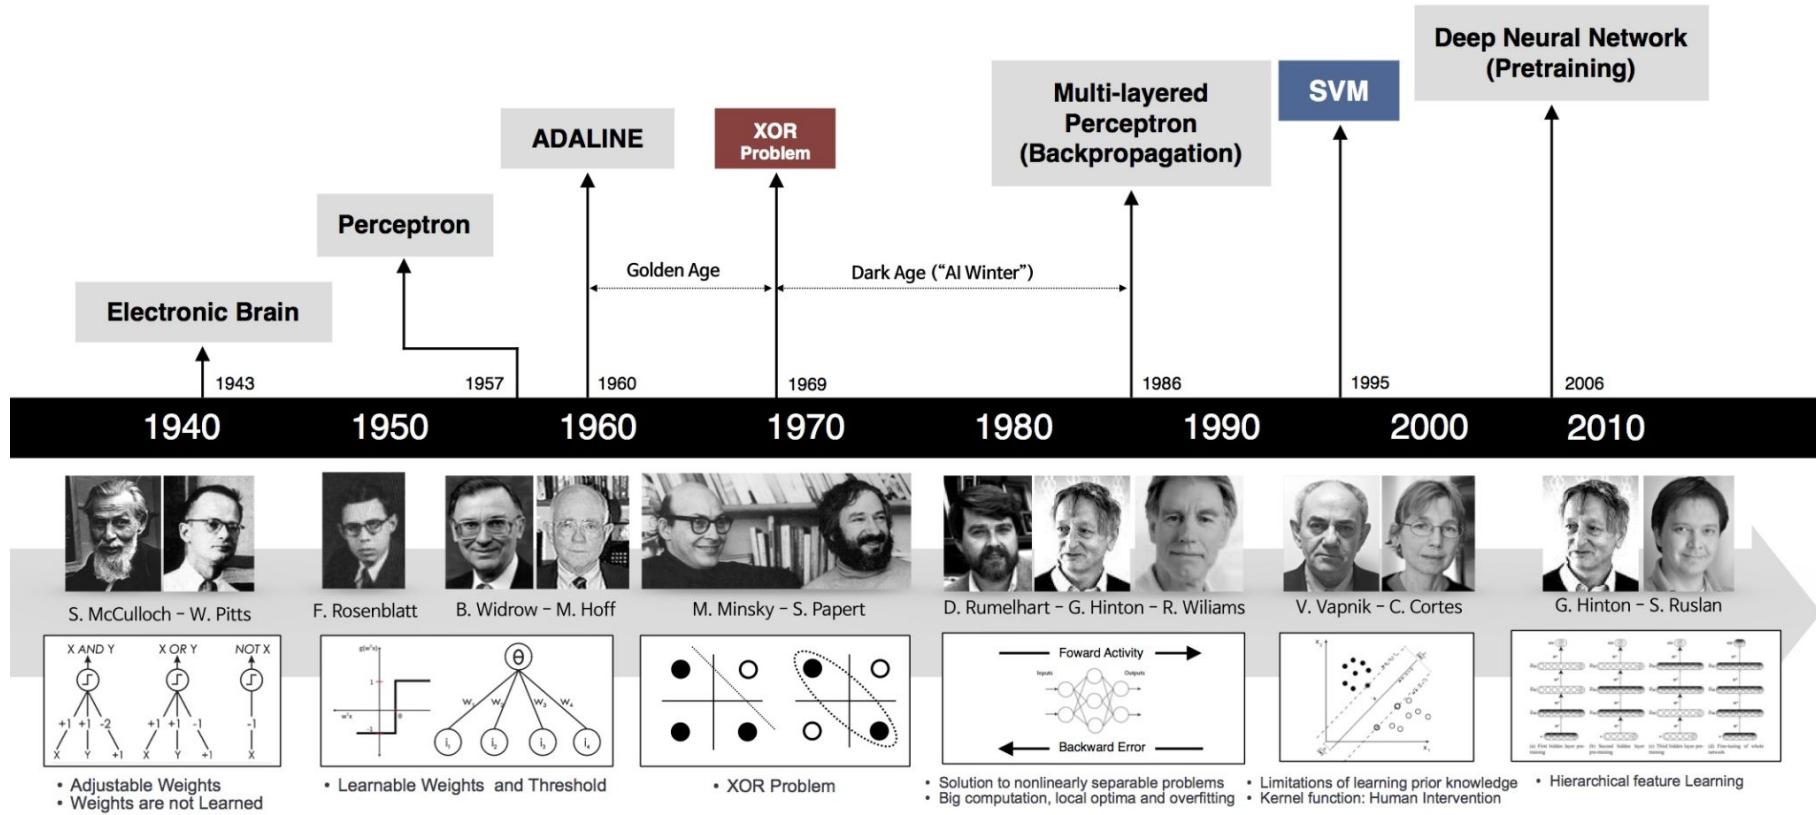
$x_1$	$x_2$	$y$
1	1	1
1	0	0
0	1	0
0	0	0

# Perceptron algorithm



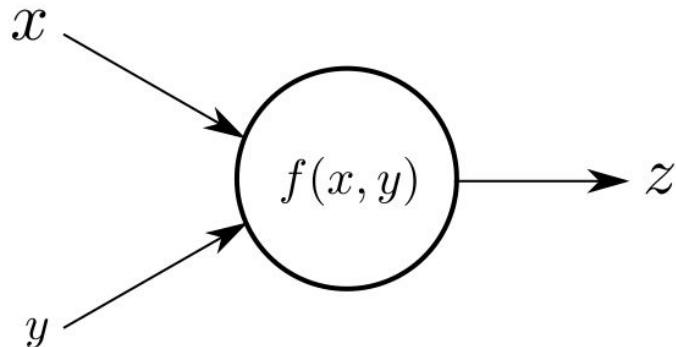
$x_1$	$x_2$	$y$
1	1	0
1	0	1
0	1	1
0	0	0

# Timeline

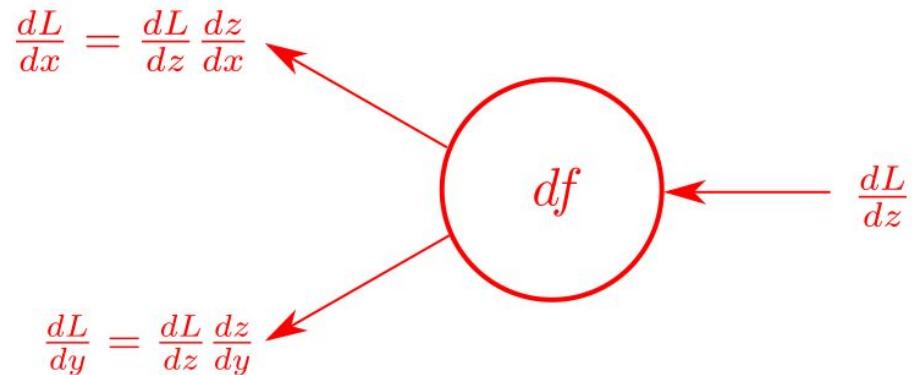


# Backprop

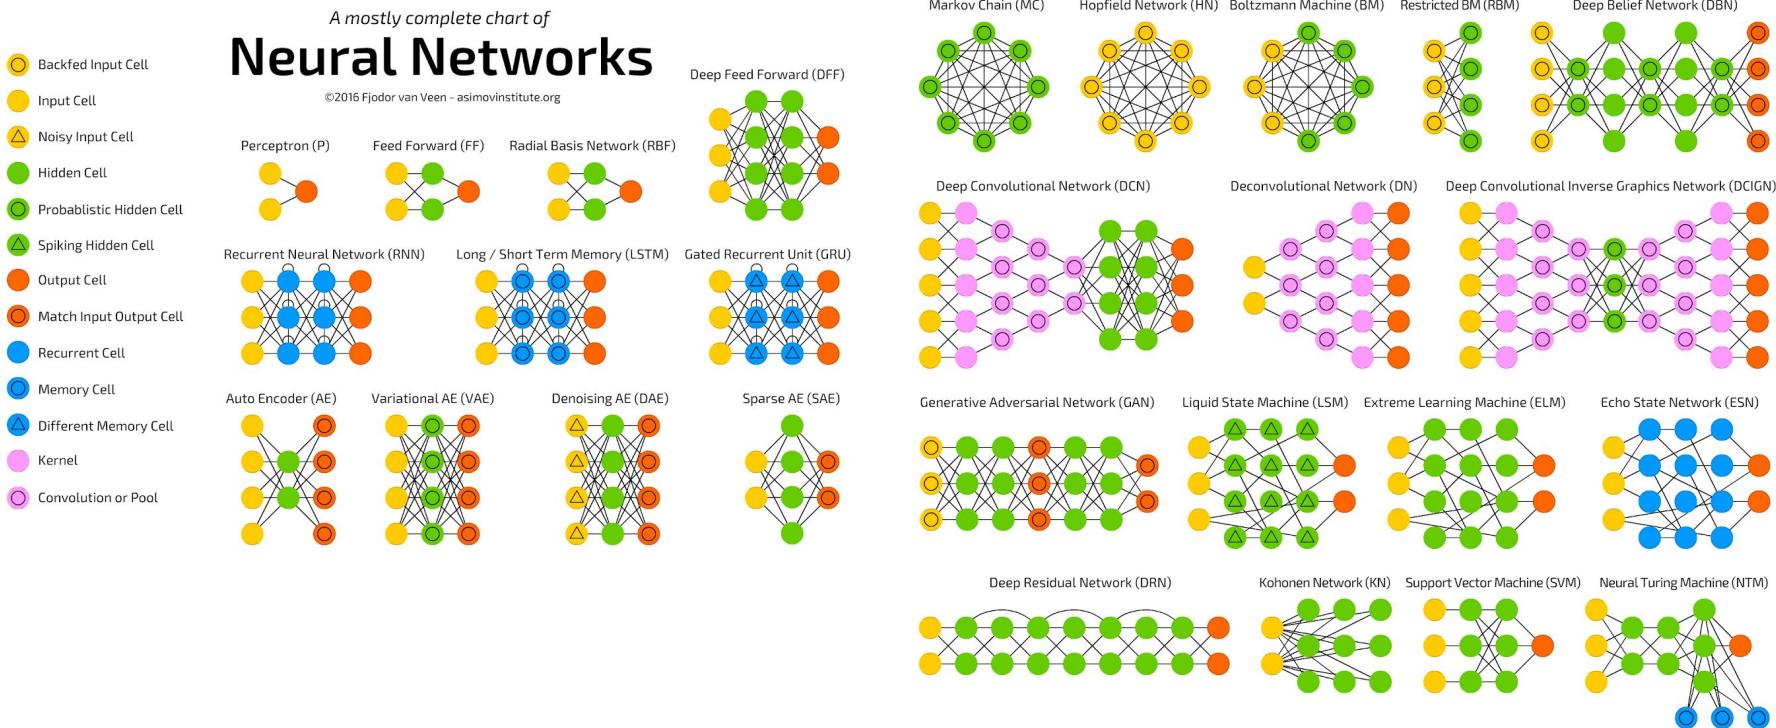
Forwardpass



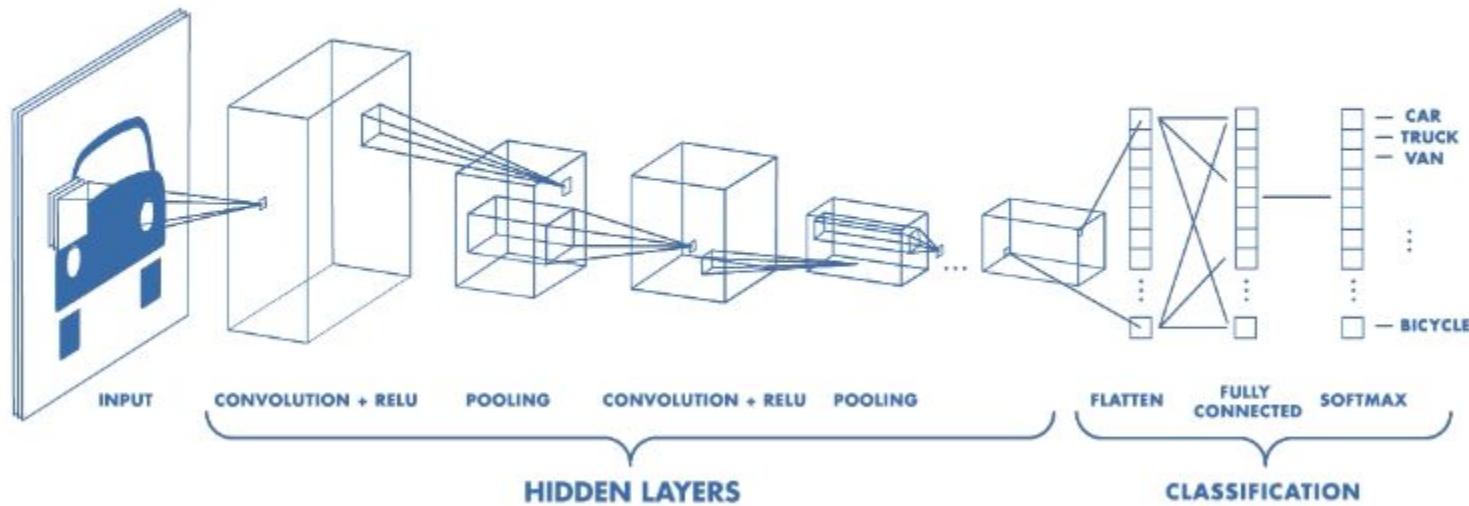
Backwardpass



# Some architectures



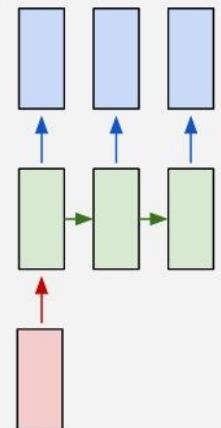
# Convolutional Neural Nets



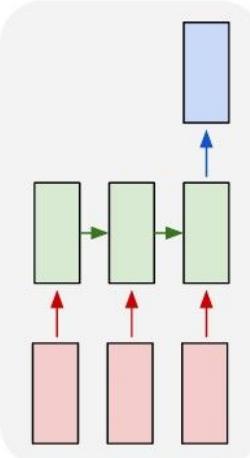
<https://towardsdatascience.com/3bd2b1164a53>

# Recurrent Neural Nets

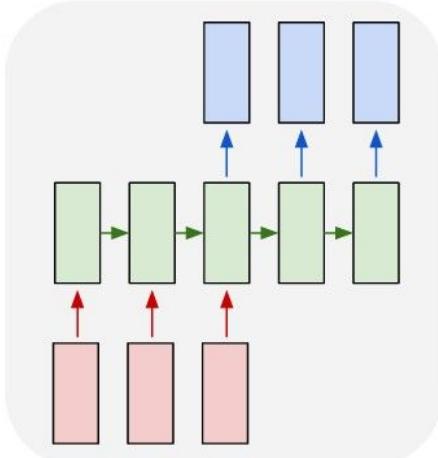
one to many



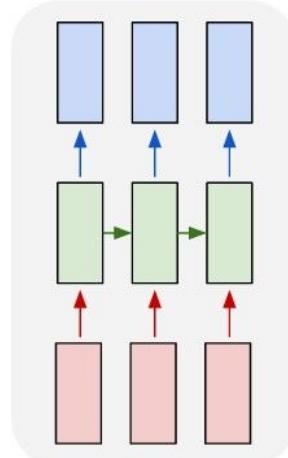
many to one



many to many



many to many



# Building blocks

Layers

Activation functions

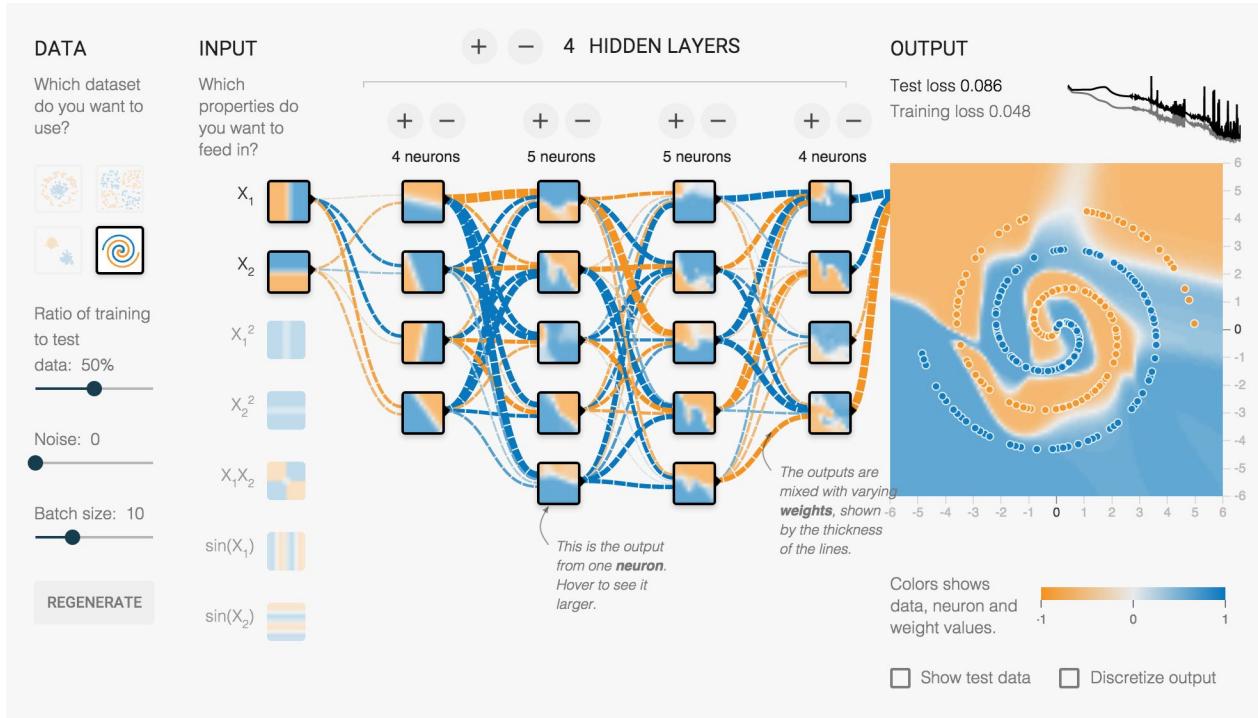
Loss functions

Optimizers

Regularizers

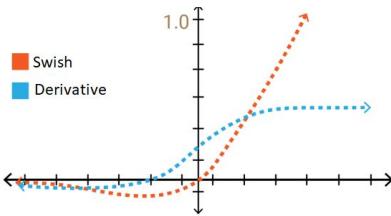
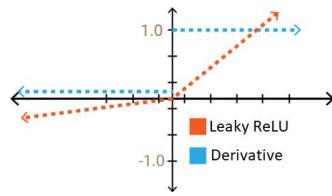
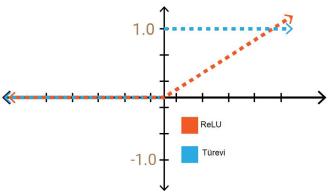
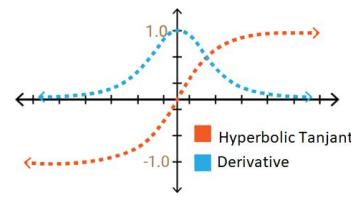
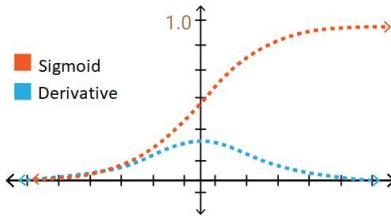
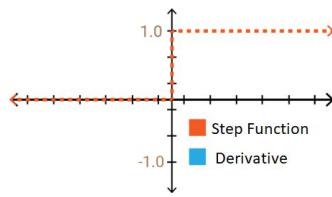
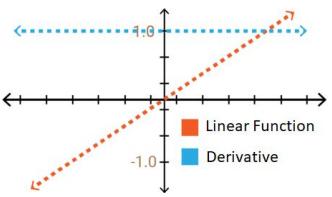
Evaluation metrics

# Layers



<https://blogs.scientificamerican.com/sa-visual/unveiling-the-hidden-layers-of-deep-learning/>

# Activation functions

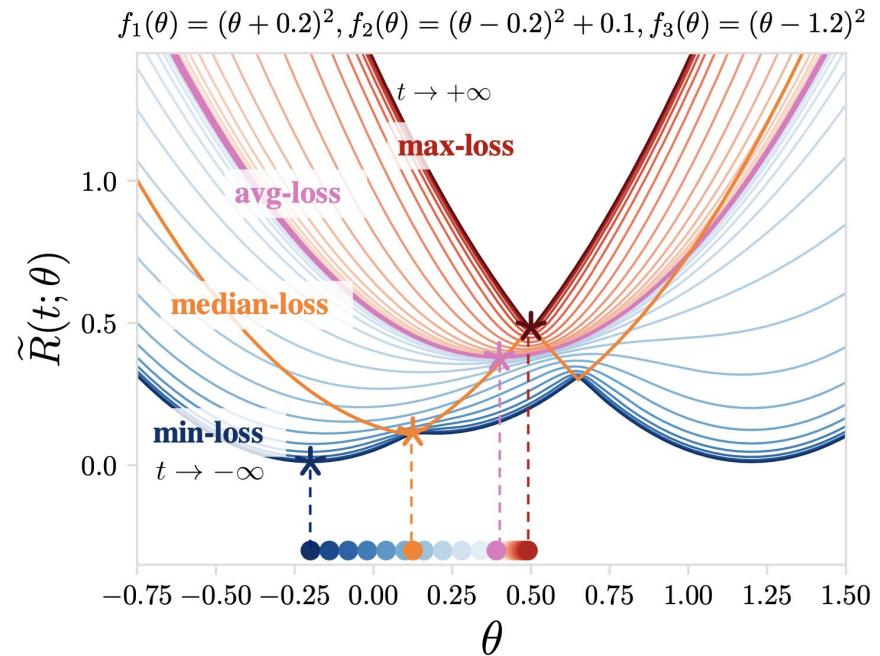


<https://towardsdatascience.com/706ac4284c8a> and [https://en.wikipedia.org/wiki/Activation\\_function](https://en.wikipedia.org/wiki/Activation_function)

# Loss functions

**Goal:** Minimize empirical loss

$$J(\mathbf{W}) = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(\hat{y}_i, y_i)$$



<https://pythonawesome.com/tilted-empirical-risk-minimization-in-python/>

# Loss functions for classification tasks

**Binary Crossentropy loss** for  $C = 2$  classes:

$$\mathcal{L}(\hat{y}_i, y_i) = -\frac{1}{N} \sum_i^N y_i \log \hat{y}_i + (1 - y_i) \log (1 - \hat{y}_i)$$

**Categorical Crossentropy loss** for  $C > 2$  classes:

$$\mathcal{L}(\hat{y}_i, y_i) = -\frac{1}{N} \sum_i^N \sum_c^C \mathbb{1}_{y_i \in C_c} \log \hat{y}_i$$

Note:  $\hat{y}_i = p(y_i)$

# Loss functions for regression tasks

Mean Squared Error loss:

$$\mathcal{L}(\hat{y}_i, y_i) = (y_i - \hat{y}_i)^2$$

Notes:

$$\text{RMSE} = \sqrt{\text{MSE}}$$

$$\text{MAE} \leq \text{RMSE}$$

Mean Absolute Error loss:

$$\mathcal{L}(\hat{y}_i, y_i) = |y_i - \hat{y}_i|$$

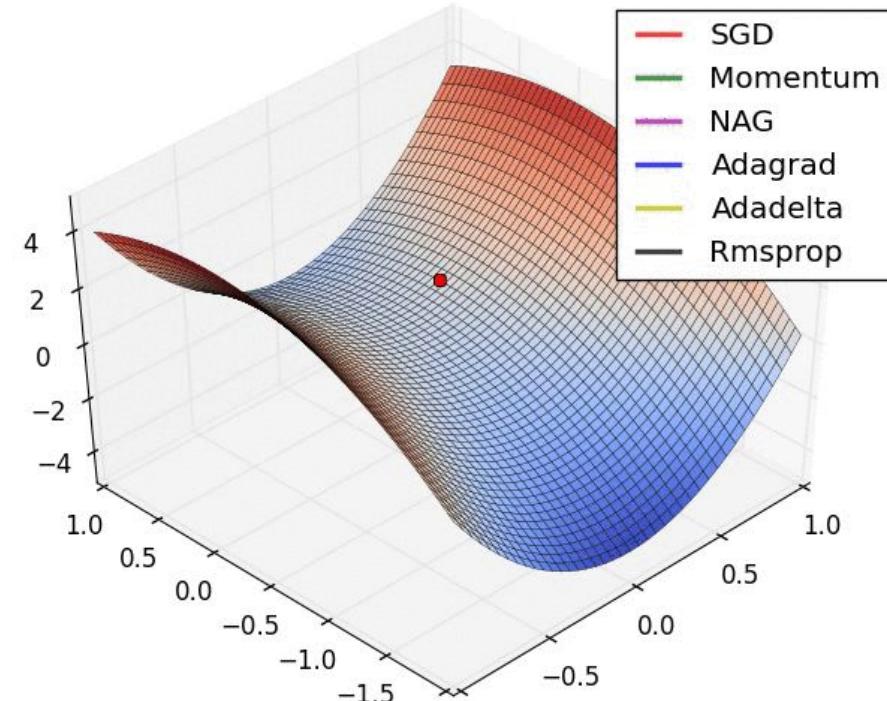
$$\text{RMSE} \leq \text{MAE} \cdot \sqrt{n}$$

# Optimizers

Mostly based on Gradient Descent:

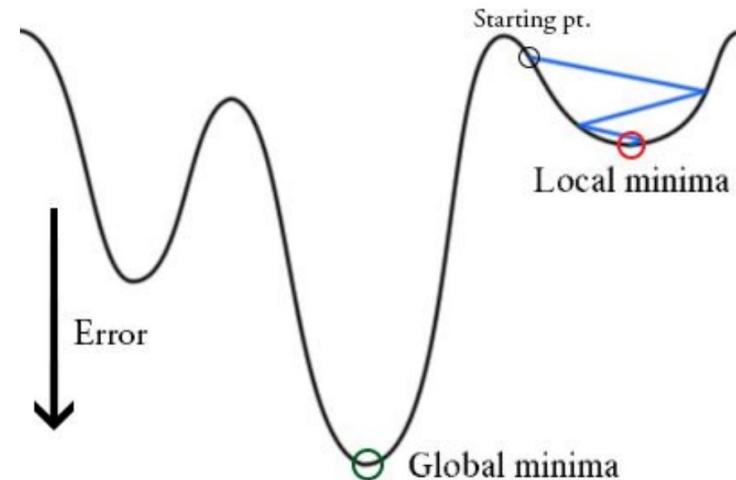
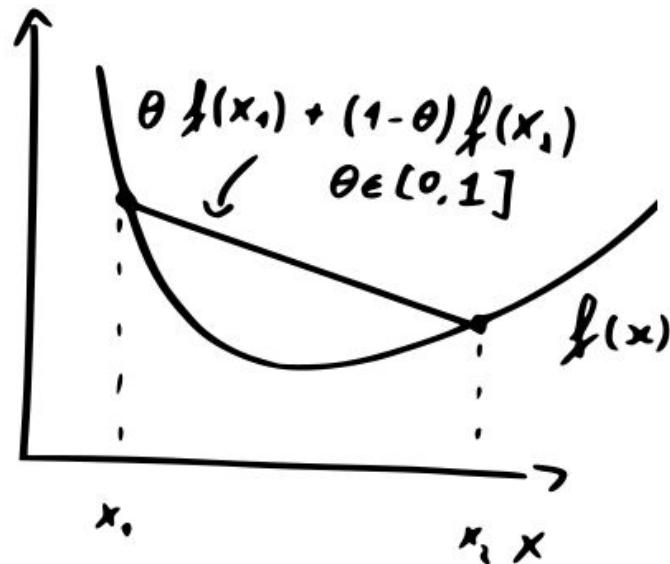
$$\theta = \theta - \eta \cdot \nabla J(\mathbf{W}; \theta)$$

- SGD
- RMSProp
- Adam
- AdaGrad
- AdaDelta



<https://ruder.io/optimizing-gradient-descent/>

# Non-convex optimization



<https://michielsstock.github.io/posts/2018-03-07-ConvexSummary/>

[https://www.cs.ubc.ca/labs/lci/mlrg/slides/non\\_convex\\_optimization.pdf](https://www.cs.ubc.ca/labs/lci/mlrg/slides/non_convex_optimization.pdf)

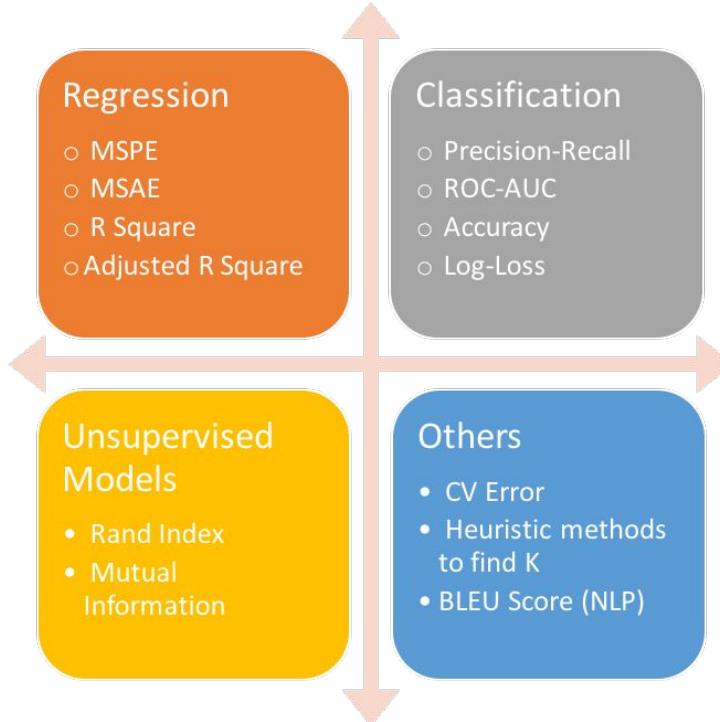
# Regularizers

**Goal:** Improve model generalization

- Dropout\*
- BatchNorm
- Early stopping\*
- L1 & L2 penalty
- Adaptive learning rate\*
- Data augmentation

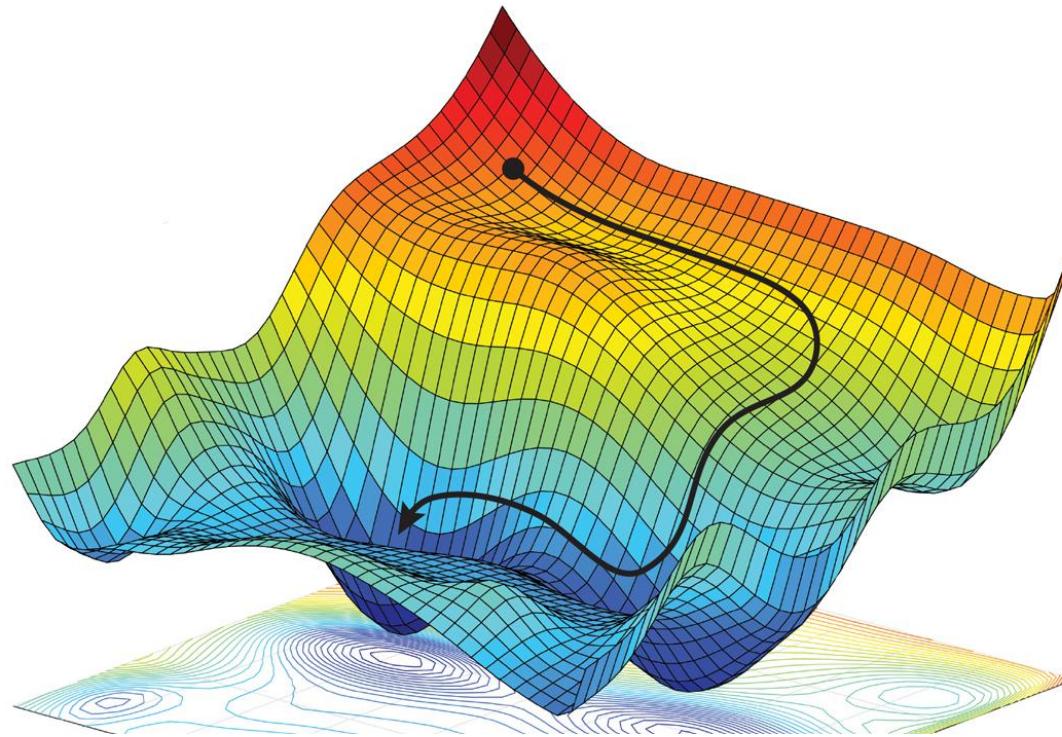
\* Very very useful, really

# Evaluation metrics



# Key notes

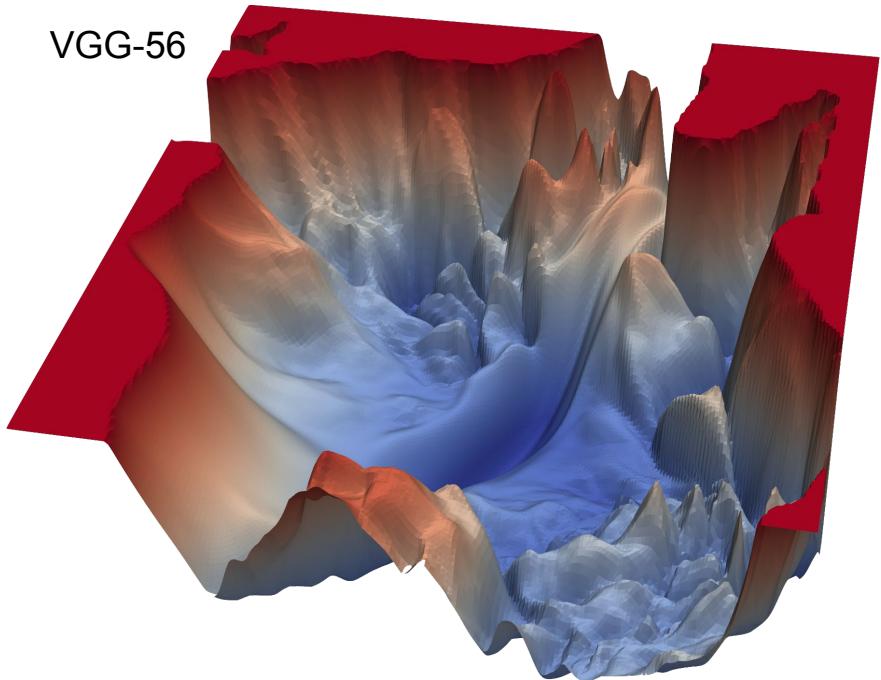
# DL in practice



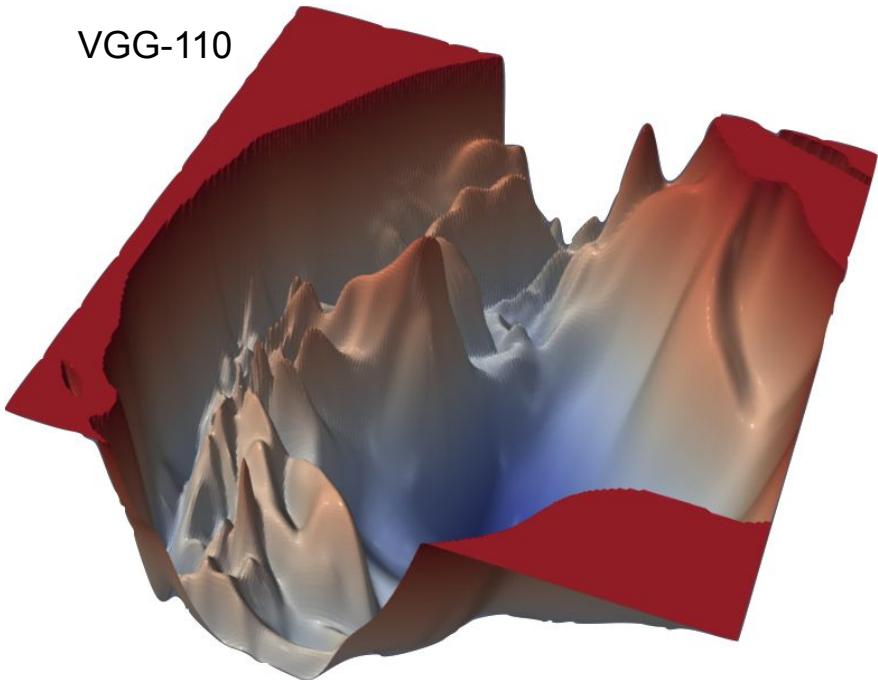
<https://www.science.org/content/article/ai-researchers-allege-machine-learning-alchemy>

# Loss landscapes

VGG-56



VGG-110



# Generalization

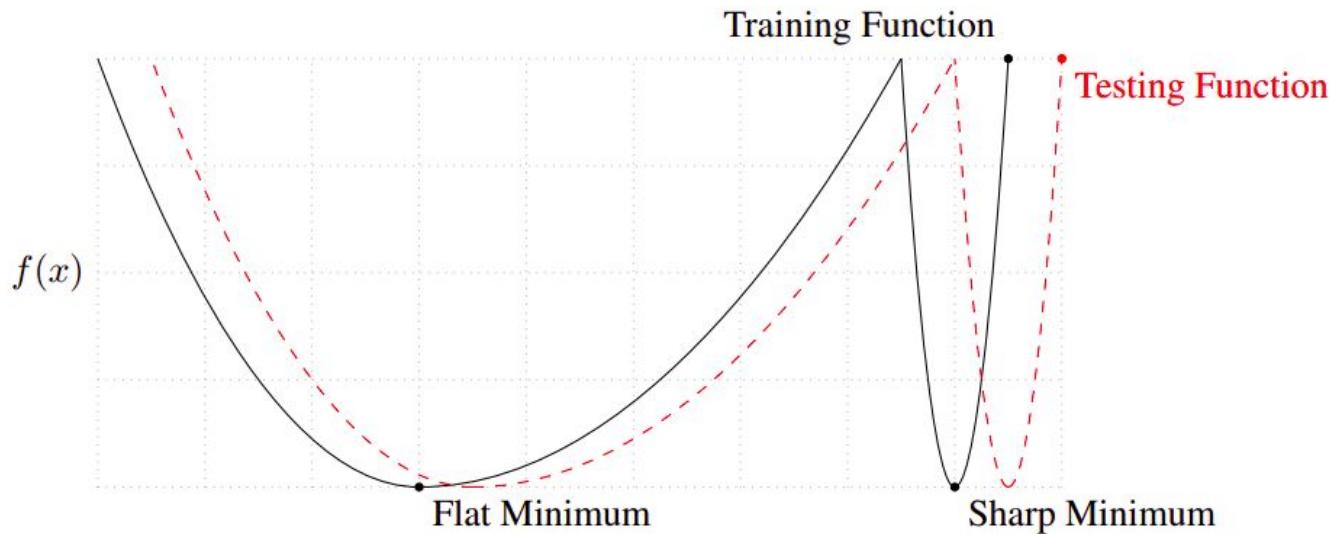
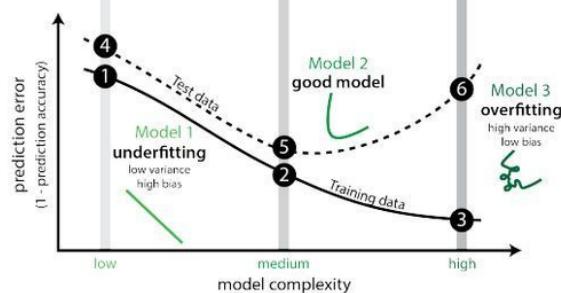
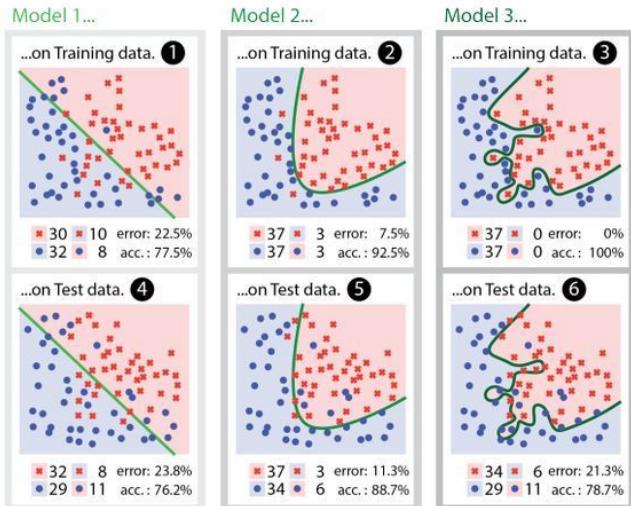
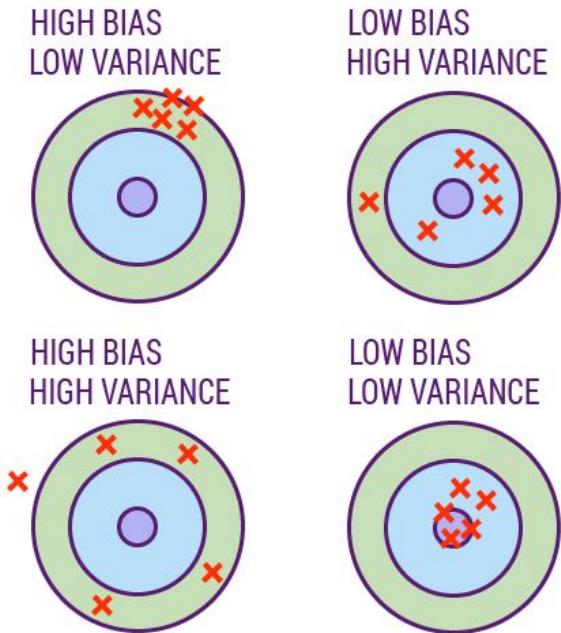
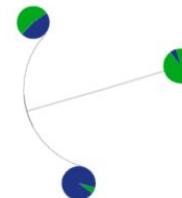
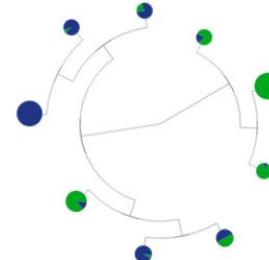
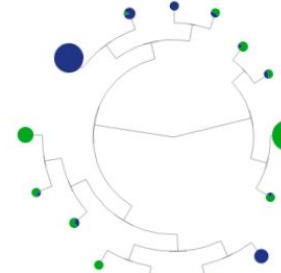
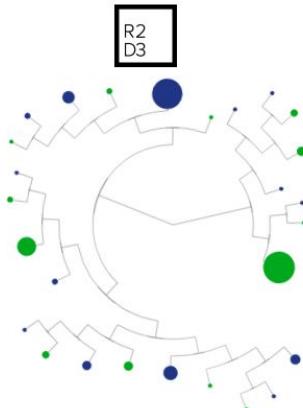


Figure 1: A Conceptual Sketch of Flat and Sharp Minima. The Y-axis indicates value of the loss function and the X-axis the variables (parameters)

# Bias-Variance tradeoff



# Bias-Variance tradeoff interactive demo

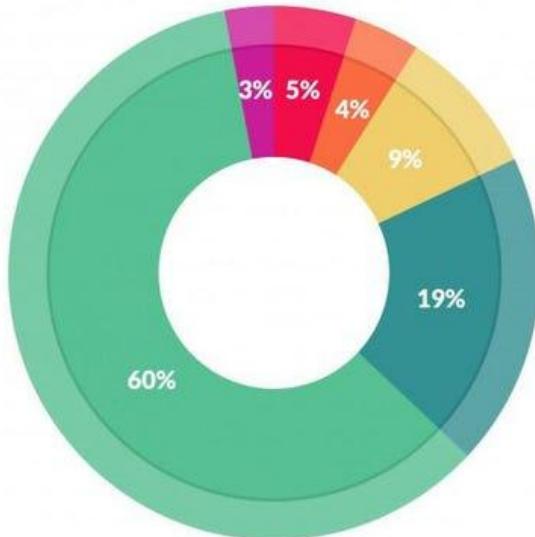


A VISUAL INTRODUCTION  
TO MACHINE LEARNING—PART II

## Model Tuning and the Bias-Variance Tradeoff

<http://www.r2d3.us/visual-intro-to-machine-learning-part-2/>

# Data science



## What data scientists spend the most time doing

- *Building training sets: 3%*
- *Cleaning and organizing data: 60%*
- *Collecting data sets; 19%*
- *Mining data for patterns: 9%*
- *Refining algorithms: 4%*
- *Other: 5%*

<https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/>

# Limitations of DL

Lack of reasoning

Lack of generalization → DL excels at *automating* tasks

Lack of prior knowledge → solvable with transfer learning

Lack of interpretability (black-box models) → XAI working on this

Lots of training data often needed → Recent research on *Efficient DL*

Models prone to adversarial noise → Hot topic!

DL is a “too-empirical” field → but so do HCI

# Some tricks and tips

<http://cs231n.github.io/neural-networks-3/>

<http://karpathy.github.io/2019/04/25/recipe/>

<https://towardsdatascience.com/3ca24c31ddc8>

<https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-deep-learning-tips-and-tricks>

# Bibliography

1. F. Chollet. **Deep Learning with Python**, 2021
2. M. Nielsen. **Neural Networks and Deep Learning**, 2015
3. I. Goodfellow, Y. Bengio, A. Courville. **Deep Learning**, 2016
4. A. Gibson, J. Patterson. **Deep Learning: A Practitioners Approach**, 2017
5. R. O. Duda, P. E. Hart, D. G. Stork. **Pattern Classification**, 2000

